

≡

🔍

{x}

🔗


📁

<>

☰

▶

⌵ Import libraries and dataset



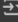
```
import matplotlib.pyplot
import pandas as pd
import numpy as np
%matplotlib inline
import seaborn as sns
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

[]

```
alphanorm = pd.read_csv('alphanorm.csv', index_col = False)
```

[]


```
alphanorm.head()
```



	sex	hb	pcv	rbc	mcv	mch	mchc	rdw	wbc	neut	lymph	plt	hba	hba2	hbf	phenotype
0	female	10.8	35.2	5.12	68.7	21.2	30.8	13.4	9.6	53.0	33.0	309.0	88.5	2.6	0.11	alpha carrier
1	male	10.8	26.6	4.28	62.1	25.3	40.8	19.8	10.3	49.4	43.1	687.0	87.8	2.4	0.90	alpha carrier
2	female	10.8	35.2	5.12	68.7	21.2	30.8	13.4	9.6	53.0	33.0	309.0	88.5	2.6	0.10	alpha carrier
3	male	14.5	43.5	5.17	84.0	28.0	33.4	12.1	11.9	31.0	50.0	334.0	86.8	2.8	0.30	alpha carrier
4	male	11.5	34.4	5.02	68.7	22.9	33.4	15.7	20.4	67.0	30.0	596.0	86.3	2.4	1.30	alpha carrier

[]

```
alphanorm.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 203 entries, 0 to 202
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   sex         203 non-null    object
1   hb          203 non-null    float64
2   pcv         203 non-null    float64
3   rbc         202 non-null    float64
4   mcv         203 non-null    float64
5   mch         201 non-null    float64
6   mchc        203 non-null    float64
7   rdw         203 non-null    float64
8   wbc         203 non-null    float64
9   neut        203 non-null    float64
10  lymph       203 non-null    float64
11  plt         203 non-null    float64
12  hba         203 non-null    float64
13  hba2        203 non-null    float64
14  hbf         203 non-null    float64
15  phenotype   203 non-null    object
dtypes: float64(14), object(2)
memory usage: 25.5+ KB
```


⌵ Handle categorical values

[]

```
alphanorm['rbc'] = alphanorm['rbc'].fillna(alphanorm.groupby('phenotype')['rbc'].transform('mean'))
alphanorm['mch'] = alphanorm['mch'].fillna(alphanorm.groupby('phenotype')['mch'].transform('mean'))
```

[]

```
for col in alphanorm.columns:
    if alphanorm[col].dtype != object:
        Q1 = alphanorm[col].quantile(0.25)
        Q3 = alphanorm[col].quantile(0.75)
        IQR = Q3 - Q1
        S = 1.5*IQR
        LB = Q1 - S
        UB = Q3 + S
        print(UB)
        alphanorm.loc[alphanorm[col] > UB,col] = UB
        alphanorm.loc[alphanorm[col] < LB,col] = LB
```



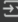
17.025000000000002
47.925000000000002
6.537500000000001
104.32500000000003
35.125
36.29074075
19.924999999999997
15.125000000000002
65.875
59.5
564.75
90.61428571249999
3.000000000000004
1.4730769225

[]

```
alphanorm = alphanorm.astype({'sex': 'category', 'phenotype': 'category'})
```

[]

```
alphanorm['phenotype'] = alphanorm['phenotype'].replace({'alpha carrier': 'alpha carrier'})
alphanorm['phenotype'] = alphanorm['phenotype'].replace({True:1, False:0})
```



```
<ipython-input-169-e8284ae63eeb>:2: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call
alphanorm['phenotype'] = alphanorm['phenotype'].replace({True:1, False:0})
```

1

▶

- ✦ Distinguish features & label

```
[ ] x = alphanorm.drop('phenotype', axis=1)
    y = alphanorm['phenotype']
```

```
[ ] categorical_vars = list(X.select_dtypes(include=['category']).columns)
categorical_vars
```

→ ['sex']

```
[ ] x.head()
```

	sex	hb	pcv	rbc	mcv	mch	mchc	rdw	wbc	neut	lymph	plt	hba	hba2	hbf
0	female	10.8	35.2	5.12	68.7	21.2	30.800000	13.4	9.600	53.000	33.0	309.00	88.5	2.6	0.11
1	male	10.8	26.6	4.28	62.1	25.3	36.290741	19.8	10.300	49.400	43.1	564.75	87.8	2.4	0.90
2	female	10.8	35.2	5.12	68.7	21.2	30.800000	13.4	9.600	53.000	33.0	309.00	88.5	2.6	0.10
3	male	14.5	43.5	5.17	84.0	28.0	33.400000	12.1	11.900	31.000	50.0	334.00	86.8	2.8	0.30
4	male	11.5	34.4	5.02	68.7	22.9	33.400000	15.7	15.125	65.875	30.0	564.75	86.3	2.4	1.30

- Perform encoding & Split the dataset

```
[ ] from sklearn.preprocessing import OneHotEncoder
    from sklearn.compose import ColumnTransformer
```

```
one_hot = OneHotEncoder()  
transformer = ColumnTransformer([('one_hot', one_hot, categorical_vars)],  
                               remainder='passthrough')  
X = pd.DataFrame(transformer.fit_transform(X))
```

```
[ ] from sklearn.model_selection import StratifiedShuffleSplit
split = StratifiedShuffleSplit(n_splits=1, test_size=0.20, random_state=42)
for train_index, test_index in split.split(alphanorm, alphanorm["phenotype"]):
    strat_train = alphanorm.loc[train_index]
    strat_test = alphanorm.loc[test_index]
```

```
[ ] train_X = strat_train.drop('phenotype', axis=1)
train_y = strat_train['phenotype']
test_X = strat_test.drop('phenotype', axis=1)
test_y = strat_test['phenotype']
```

```
[ ] one_hot1 = OneHotEncoder()
transformer = ColumnTransformer([(['one_hot', one_hot1, categorical_vars],
                                remainder= 'passthrough')]
train_X = transformer.fit_transform(train_X)
train_X = pd.DataFrame(train_X)
```

```
[ ] one_hot2 = OneHotEncoder()
transformer = ColumnTransformer([('one_hot', one_hot2, categorical_vars)],
                                remainder='passthrough')
test_x = transformer.fit_transform(test_x)
test_x = pd.DataFrame(test_x)
```

```
[ ] train_X.head()
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1.0	0.0	11.3	35.900000	4.71	76.0	24.0	31.5	14.0	8.90	65.875	29.0	268.0	86.523291	2.588608	0.769231
1	1.0	0.0	11.6	36.000000	4.69	77.0	25.0	32.2	13.3	7.70	51.000	59.5	249.0	86.600000	2.500000	0.100000
2	1.0	0.0	13.8	35.876404	5.81	74.0	23.7	31.6	11.8	9.90	53.000	37.0	312.0	87.100000	2.600000	0.100000
3	0.0	1.0	11.8	36.000000	5.27	68.3	22.4	32.5	14.8	12.48	42.000	51.0	311.0	86.600000	2.700000	0.500000
4	0.0	1.0	13.4	42.600000	5.13	83.0	26.2	31.6	12.0	7.10	43.000	50.0	235.0	84.864586	2.679310	0.537931

- ✓ Perform normalization & over sampling

```
[ ] from sklearn import preprocessing
train_x = preprocessing.normalize(train_x)
test_x = preprocessing.normalize(test_x)
```

```
[ ] from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state = 2)
train_X, train_y = sm.fit_resample(train_X, train_y)
```

```
[ ] train_X
```

```
array([[3.26934925e-03, 0.00000000e+00, 3.69436466e-02, ...,
        8.28748575e-01, 8.46306231e-03, 2.51488400e-03],
       [3.43141227e-03, 0.00000000e+00, 3.98043823e-02, ...,
        2.97160302e-01, 8.57853067e-03, 3.34141272e-04],
       [2.91233460e-03, 0.00000000e+00, 4.01902174e-02, ...,
        2.53664343e-01, 7.57206995e-03, 2.91233460e-04],
       ...,
       [1.15035313e-04, 3.08944499e-03, 3.59361935e-02, ...,
        2.77420552e-01, 8.29645257e-03, 2.38799967e-03],
       [0.00000000e+00, 3.22289376e-03, 4.35363515e-02, ...,
        2.75678880e-01, 7.98737908e-03, 1.18190454e-03],
       [2.8642670e-03, 3.27939116e-04, 4.44285301e-02, ...,
```

```
[ ] train_y.value_counts()
```

Random Forest Classifier

```
[ ] # rf = RandomForestClassifier(n_estimators=50, criterion='gini', max_features='sqrt')
    rf = RandomForestClassifier(n_estimators=100, criterion='entropy', max_features='log2', random_state=0)
```

```
RandomForestClassifier(criterion='entropy', max_features='log2', random_state=0)
```

```
[ ] y_pred
```

```
[ ] print("Accuracy:", accuracy_score(test_y, y_pred))
print("Confusion Matrix:\n", confusion_matrix(test_y, y_pred))
print("Classification Report:\n", classification_report(test_y, y_pred))
```

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted sample
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))

```

✓ Bagging Classifier

```
[ ] bagging = BaggingClassifier(estimator=RandomForestClassifier(criterion='gini', max_features='sqrt'), n_estimators=50, random_state=0)
```

```

> BaggingClassifier
  estimator:
    RandomForestClassifier

```

```
[ ] accuracy_score(test_y, y_pred)
```

AdaBoost Classifier

```
[ ] ada.fit(train_X, train_y)
```

```
AdaBoostClassifier
AdaBoostClassifier(n_estimators=100, random_state=0)
```

```
[ ] y_pred = ada.predict(test_x)
```

```
[ ] accuracy_score(test_y, y_pred)
```

```
0.6585365853658537
```

▼ KNN Classifier

```
[ ] from sklearn.neighbors import KNeighborsClassifier
```

```
[ ] knn = KNeighborsClassifier(n_neighbors = 10, metric='minkowski', p=5)
```

```
[ ] knn.fit(train_X, train_y)
```

```
KNeighborsClassifier
KNeighborsClassifier(n_neighbors=10, p=5)
```

```
print("accuracy score: ", knn.score(test_x, test_y))
print("confusion matrix: \n", confusion_matrix(test_y, knn.predict(test_x)))
print("classification report: ", classification_report(test_y, knn.predict(test_x)))
```

```
accuracy score: 0.4146341463414634
confusion matrix:
[[ 7  4]
 [20 10]]
classification report:

```

			precision	recall	f1-score	support
	0	0.26	0.64	0.37	11	
	1	0.71	0.33	0.45	30	
accuracy			0.41	0.41		
macro avg	0.49	0.48	0.41	0.41		
weighted avg	0.59	0.41	0.43	0.41		

▼ Use Cross Validation to find best hyperparameter for RF

```
[ ] from sklearn.model_selection import cross_val_score
    from sklearn.model_selection import KFold
```

```
[ ] k = KFold(n_splits=10)
```

```
[ ] # Implement cross validation for diabetes data set
    for i, j in k.split(X):
        # Use .iloc to select rows based on indices
        X_train_cv, X_test_cv = X.iloc[i], X.iloc[j]
        y_train_cv, y_test_cv = y.iloc[i], y.iloc[j]
        rf = RandomForestClassifier(n_estimators=100, criterion='entropy', max_features='log2', random_state=0)
        rf.fit(X_train_cv, y_train_cv)
```

```
[ ] print("cross_val_score: ", cross_val_score(rf, train_X, train_y, cv = 10))
    print("\n")
    print("mean cross_val_score: ", np.mean(cross_val_score(rf, train_X, train_y, cv = 10)))
```

```
cross_val_score: [0.75      0.83333333 0.58333333 0.66666667 0.875      0.83333333
 0.91304348 0.86956522 0.86956522 0.86956522]
```

```
mean cross_val_score: 0.806340579710145
```

▼ Define accuracies of models

```
[ ] rf_accuracy = cross_val_score(estimator=rf, X=train_X, y=train_y, cv=10)
    knn_accuracy = cross_val_score(estimator=knn, X=train_X, y=train_y, cv=10)
    ada_accuracy = cross_val_score(estimator=ada, X=train_X, y=train_y, cv=10)
    bagging_accuracy = cross_val_score(estimator=bagging, X=train_X, y=train_y, cv=10)
```

```
[ ] print("RF Accuracy: {:.2f}%".format(rf_accuracy.mean()*100))
    print("Ada Boost Accuracy: {:.2f}%".format(ada_accuracy.mean()*100))
    print("KNN Accuracy: {:.2f}%".format(knn_accuracy.mean()*100))
    print("Bagging Accuracy: {:.2f}%".format(bagging_accuracy.mean()*100))
```

```
RF Accuracy: 80.63%
Ada Boost Accuracy: 73.08%
KNN Accuracy: 57.16%
```

▼ Use GridSearchCV to find the best hyperparameters

```
[ ] from sklearn.model_selection import GridSearchCV
```

▼ For RF Model

```
[ ] param_grid = {
    'n_estimators': [10, 50, 100],
    'max_features': ['auto', 'sqrt', 'log2'],
    'criterion': ['gini', 'entropy']
}
```

```
[ ] grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=10, scoring='accuracy', n_jobs=-1)
grid_search.fit(train_X, train_y)
```

Show hidden output

```
[ ] print("Best Hyperparameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
```

Best Hyperparameters: {'criterion': 'entropy', 'max_features': 'log2', 'n_estimators': 100}
Best Accuracy: 0.8320652173913043

▼ For KNN Model

```
[ ] param_grid = {
    'n_neighbors': [3, 5, 7, 9],
    'weights': ['uniform', 'distance'],
    'metric': ['minkowski', 'euclidean']
}
```

```
[ ] grid_search = GridSearchCV(estimator=knn, param_grid=param_grid, cv=10, scoring='accuracy', n_jobs=-1)
grid_search.fit(train_X, train_y)
```

GridSearchCV

best_estimator_: KNeighborsClassifier

KNeighborsClassifier

```
print("Best Hyperparameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
```

Best Hyperparameters: {'metric': 'minkowski', 'n_neighbors': 3, 'weights': 'distance'}
Best Accuracy: 0.6735507246376812

▼ For Bagging Model

```
[ ] param_grid = {
    'n_estimators': [10, 50, 100],
    'max_samples': [0.5, 0.75, 1.0]
}
```

```
[ ] grid_search = GridSearchCV(estimator=bagging, param_grid=param_grid, cv=10, scoring='accuracy', n_jobs=-1)
grid_search.fit(train_X, train_y)
```

```
[ ] print("Best Hyperparameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
```

▼ For AdaBoost Model

```
[ ] param_grid = {
    'n_estimators': [10, 50, 100],
    'learning_rate': [0.1, 0.5, 1.0]
}
```

Best Hyperparameters: {'learning_rate': 1.0, 'n_estimators': 100}
Best Accuracy: 0.7307971014492753

```
[ ] grid_search = GridSearchCV(estimator=ada, param_grid=param_grid, cv=10, scoring='accuracy', n_jobs=-1)
grid_search.fit(train_X, train_y)
```

```
[ ] print("Best Hyperparameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
```

▼ Create pickle

```
[ ] import joblib
# Save the model as a pickle in a file
joblib.dump(rf, 'thalassemia_model.pkl')
```

['thalassemia_model.pkl']



[Colab paid products](#) - [Cancel contracts here](#)

✓ Connected to Python 3 Google Compute Engine backend

